

# MAD Adv I

# About Me

- Maruthi R Janardhan
  - Been with IBM, ANZ, HCL-HP, my own startup Leviossa..
  - Total 16 years programming/consulting C, C++, Java, Javascript, Ruby, Perl, Python, PHP, etc
  - Two pending Indian Patents

# Polyglot Databases

- Understand Consistency Models - ACID/BASE
- Whats driving us to this
- Benefits
- Challenges
- Deciding Criteria

# ACID/BASE

- Limits of instantaneous consistency
  - The balance between write performance and availability in instantaneous consistent systems
- What does Eventual Consistency get us
- CAP Theorem

# Deciding Criteria

- Write performance
- Read performance
- Availability
- Consistency
- Partition Tolerance
- Query patterns
  - Text search, graph search

# Types of DBs

- Mongo
  - Tuneable consistency, Availability, Partition Tolerance.
  - Indexed Document Store
  - Master slave model with writes only to master - Easy scalability for reads
- Neo4j
  - Graph database tuned for graph questions
  - Consistent and Available. Billions of nodes and relationships

# Types of DB

- Cassandra
  - Column store with multi master ability. Closest to SQL
  - Available and Partition Tolerant
  - Easy scalability for writes
- Redis in-memory key-value store

# Question of Durability

- What happens to backups
- How facebook handles redundancy and durability



# Threading and Locking

- Understand thread racing
- Solving thread racing with locks
- Using atomic variables

# Shared Collections

- Using shared collections
  - Write and read operations. Try multiple thread read and write into a synchronized map
  - Compare it with concurrent map implementation
- Fail fast and snapshot iterators
- CopyOnWriteArrayLists

# I/O Performance

- Java IO library read operation
- Asynchronous Read using NIO
- Measure performance in multi threaded environment

# Async Programming Model

- NodeJs
  - Demo of code
- Create a web project and perform 3 chained AJAX calls
- Performance of Async vs Sync code

# Node.js

- Node.js is server side framework for javascript that runs on google's chrome V8 engines.
- Node.js is single threaded event based server framework
- Comes with a lot of functionality “compiled” and bundled into the runtime beyond the basic JS spec

# Modules

- All of Node's functionality is bundled into modules and we have to require the modules to be able to use them.
- “fs” module contains a lot of operations related to the filesystem.
- Using the below functions, write a function that gets a list of files in a directory and another to make directory (<http://nodejs.org/api/fs.html>)

```
var fs = require("fs");  
fs.existsSync(path);  
fs.statSync(path);  
fs.readdirSync(path);  
fs.mkdirSync(newDir);
```

```
function getFilesInDir(path, callback){
  fs.exists(path, function(exists){
    if(exists){
      console.log("Checking if directory...");
      function statFunction(err, stat){
        if(err!==null){
          console.log("Error");
          callback(err);
          return;
        }
        if(!stat.isDirectory()){
          console.log("Is not a directory");
          callback("Is not a directory");
        }
        console.log("Reading directory...");
        fs.readdir(path, function(err, filesArr){
          if(err!==null){
            console.log("Error");
            callback(err);
            return;
          }
          console.log(filesArr);
          callback(filesArr);
        });
      }
      fs.stat(path, statFunction);
    }else{
      console.log("Path does not exist");
      callback("Path does not exist");
    }
  });
}
```

- Create a simple dynamic web project with 3 hardcoded json
- Create sequential AJAX calls using jquery



# Vert.x

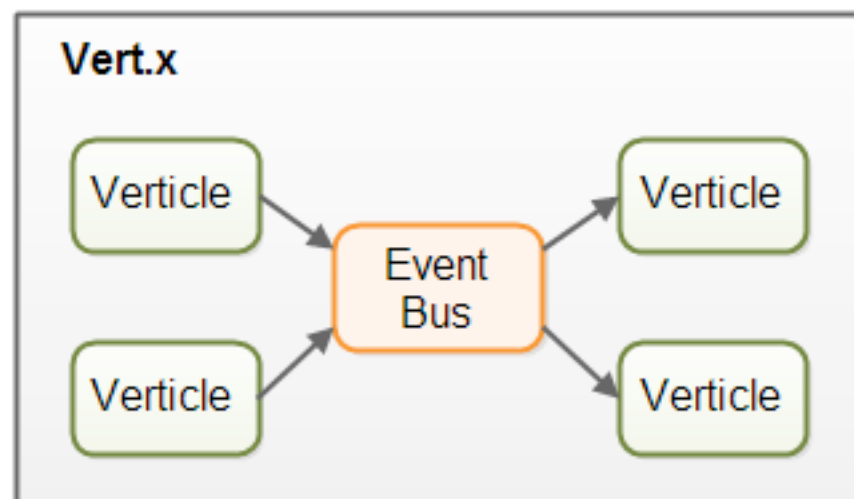
- We can build reactive, non blocking, event driven apps
- Polyglot - javascript, java, groovy and ruby
- Not an app server, modular
- Good fit for microservices

# Vert.x Performance

Best plaintext responses per second, i7-2600K hardware (55 tests)										
Framework	Best performance (higher is better)		Cls	Lng	Plt	FE	Aos	IA	Errors	
vertx	656,119	<div><div></div></div> 100.0%	Plt	Jav	vtx	Non	Lin	Rea	0	
netty	632,596	<div><div></div></div> 96.4%	Plt	Jav	Nty	Non	Lin	Rea	10	
undertow	614,547	<div><div></div></div> 93.7%	Plt	Jav	Utw	Non	Lin	Rea	70	
plain-windows	611,095	<div><div></div></div> 93.1%	Ful	Sca	Pla	Non	Win	Rea	52	
plain	543,670	<div><div></div></div> 82.9%	Ful	Sca	Pla	Non	Lin	Rea	21,473	
cpoll_cppsp	522,423	<div><div></div></div> 79.6%	Mcr	C++	Cpl	Non	Lin	Rea	0	
plain-servlet-linux	450,462	<div><div></div></div> 68.7%	Ful	Sca	Pla	Non	Lin	Rea	3,406	
jetty-servlet	448,947	<div><div></div></div> 68.4%	Plt	Jav	Jty	Jty	Lin	Rea	351	
grizzly	425,172	<div><div></div></div> 64.8%	Mcr	Jav	Svt	Grz	Lin	Rea	149	
gemini	424,586	<div><div></div></div> 64.7%	Ful	Jav	Svt	Non	Lin	Rea	2,071	
spray	422,431	<div><div></div></div> 64.4%	Mcr	Sca	Spr	Non	Lin	Rea	1,003	
servlet	417,619	<div><div></div></div> 63.6%	Plt	Jav	Svt	Res	Lin	Rea	1,214	
go	367,274	<div><div></div></div> 56.0%	Plt	Go	Go	Non	Lin	Rea	230	
openresty	330,829	<div><div></div></div> 50.4%	Plt	Lua	OpR	ngx	Lin	Rea	1,517	
spark	226,365	<div><div></div></div> 34.5%	Mcr	Jav	Svt	Res	Lin	Rea	0	
revel	221,603	<div><div></div></div> 33.8%	Ful	Go	Go	Non	Lin	Rea	521	
falcore	172,827	<div><div></div></div> 26.3%	Mcr	Go	Go	Non	Lin	Rea	0	
compojure	127,671	<div><div></div></div> 19.5%	Mcr	Clj	Svt	Res	Lin	Rea	0	
http-listener	123,434	<div><div></div></div> 18.8%	Plt	C#	Net	hts	Win	Rea	0	
falcon-pypy	112,550	<div><div></div></div> 17.2%	Mcr	Py	Gun	Non	Lin	Rea	0	
wsgi-nginx-uwsgi	110,522	<div><div></div></div> 16.8%	Plt	Py	uWS	ngx	Lin	Rea	459,212	
evhttp-sharp	107,494	<div><div></div></div> 16.4%	Mcr	C#	Mon	Non	Lin	Rea	0	
bottle-pypy	88,328	<div><div></div></div> 13.5%	Mcr	Py	Tor	Non	Lin	Rea	52	
nodejs	80,363	<div><div></div></div> 12.2%	Plt	JS	njs	Non	Lin	Rea	0	

# Verticles

- Vert.x can deploy and execute components called Verticles.
- You can think of verticles as being similar to servlets or message driven EJBs driven by an event bus




# Verticle Messaging

- Messages can be simple objects (e.g. Java objects), strings, CSV, JSON, binary data or whatever else you need
- Verticles can send and listen to addresses. An address is like a named channel.
- When a message is sent to a given address, all verticles that listen on that address receive the message.
- Verticles can subscribe and unsubscribe to addresses without the senders knowing.

# Threading Model

- verticle is only ever executed by a single thread, and always by the same thread.
- A single thread can distribute messages to multiple verticles.
- Vert.x creates one thread per CPU
- Vert.x comes with a set of built-in services (functionality).  
Some of these services are:
  - HTTP server
  - JDBC connector
  - MongoDB connector



# Number of Instances

- Deploy multiple verticle instances to run in different threads
- `DeploymentOptions options = new DeploymentOptions().setInstances(16);`

# Creating Verticles

```
public class MyVerticle extends AbstractVerticle {  
  
    @Override  
    public void start(Future<Void> startFuture) {  
        System.out.println("MyVerticle started!");  
    }  
  
    @Override  
    public void stop(Future stopFuture) throws Exception {  
        System.out.println("MyVerticle stopped!");  
    }  
  
}  
  
public static void main(String[] args) {  
    VertxOptions options = new VertxOptions().setWorkerPoolSize(10);  
    Vertx vertx = Vertx.vertx(options);  
    vertx.deployVerticle("com.mydomain.MyVerticle");  
}
```

# Verticle Events

- The start() method : start HTTP or TCP server, register event handlers on the event bus, deploy other verticles, or whatever else your verticle needs to do its work.

- Shutdown stuff in stop method

- The verticle will be deployed asynchronously

```
vertx.deployVerticle("com.mydomain.MyVerticle", new  
Handler<AsyncResult<String>>() {  
    @Override  
    public void handle(AsyncResult<String> stringAsyncResult) {  
        System.out.println("Verticle deployment complete");  
    }  
});
```



# Registering to Events

- When a verticle wants to listen for messages from the event bus, it listens on a certain address. An address is just a name (a String) which you can choose freely.
- An address is thus more like the name of a channel with multiple receivers

```
vertx.eventBus().consumer("Channel1", message -> {  
    System.out.println("message.body() = "  
        + message.body());  
});  
vertx.eventBus().publish("Channel1", "message 2");
```

# Types Of Verticles

- Standard Verticles
  - These are the most common and useful type - they are always executed using an event loop thread.
- Worker Verticles
  - These run using a thread from the worker pool. An instance is never executed concurrently by more than one thread.

```
DeploymentOptions options = new DeploymentOptions().setWorker(true);  
vertx.deployVerticle("com.mycompany.MyOrderProcessorVerticle", options);
```

- Multi-threaded worker verticles
  - These run using a thread from the worker pool. An instance can be executed concurrently by more than one thread.

# Vert.x buffers

- Carry Binary Information in buffers. Dynamically resizable

- Can be used as message payloads

```
byte[] initialData = new byte[]{1, 2, 3};
```

```
Buffer buffer = Buffer.buffer(initialData);
```

```
buffer.setShort ( 10, (short) 127);
```

```
buffer.appendByte ((byte) 127);
```

# Running Blocking Code

- When we HAVE to invoke synchronous APIs, vertx provides a way to do that (Executes using a thread from worker pool):

```
vertx.executeBlocking(future -> {  
    // Call some blocking API that takes a significant amount of time  
    to return  
    String result = someAPI.blockingMethod("hello");  
    future.complete(result);  
}, res -> {  
    System.out.println("The result is: " + res.result());  
});
```

# Vertx Web

- Vertx is bundled with a router

```
HttpServer server = vertx.createHttpServer();  
Router router = Router.router(vertx);  
router.get("/services/users/:id").handler(new UserLoader());  
server.requestHandler(router::accept).listen(8080);
```

- In the handler

```
String id = routingContext.request().getParam("id");  
HttpServerResponse response = routingContext.response();  
response.putHeader("content-type", "application/json");  
response.end(jsonresponse)
```

# Scripting Integration

- JSR 223 allows for different scripting languages to be integrated into Java
  - javascript (nashhorn engine)
  - python (jython interpreter)
  - ruby (jruby)

# Data Type Mapping

- Usually some data type mapping is defined. Here is a def for Jython

Java Type	Python Type
char	String(length of 1)
boolean	Integer(true = not zero)
byte, short, int, long	Integer
java.lang.String, byte[], char[]	String
java.lang.Class	JavaClass
Foo[]	Array(containing objects of class or subclass of Foo)
java.lang.Object	String
orb.python.core.PyObject	Unchanged
Foo	JavaInstance representing Java class Foo

# Concise Syntax of Python

```
// print the integers from 1 to 9
for (int i = 1; i < 10; i++)
{
    System.out.println(i);
}
```

```
print the integers from 1 to 9
for i in range(1,10):
    print i
```

```
String file_name="";
try(FileReader fis = new FileReader(file_name)){
    LineNumberReader lnr = new LineNumberReader(fis);
    String line="";
    Map<String,String> phoneData = new HashMap<>();
    while((line = lnr.readLine())!=null) {
        phoneData.put(line.split(",")[0], line);
    }
    phoneData.get(name)
}
```

```
with open(file_name) as phone_book:
    book = {r.split(",")[0]: r for r in phone_book }

ret_val = book[lookup_name]
```



# Vertx on docker

- Create a Dockerfile
- `docker build --tag restvertx .`
- `docker run -p 8080:8080 restvertx &`
- Access docker host with IP: <http://192.168.99.100:8080/api/user/2>
- `docker stop restvertx`
- `docker rm $(docker ps -a -q)`

# Dockerfile

```
# Extend vert.x image
FROM vertx/vertx3

#
ENV VERTICLE_NAME com.mydomain.myapp.RestAppVerticle
ENV VERTICLE_FILE target/VertxREST-0.0.1-SNAPSHOT.jar

# Set the location of the verticles
ENV VERTICLE_HOME /usr/verticles

EXPOSE 8080

# Copy your verticle to the container
COPY $VERTICLE_FILE $VERTICLE_HOME/

# Launch the verticle
WORKDIR $VERTICLE_HOME
ENTRYPOINT ["sh", "-c"]
CMD ["vertx run $VERTICLE_NAME -cp $VERTICLE_HOME/*"]
```

# Vertx Web Handling User Data

```
router.post("/services/users").handler(new UserPersister());
```

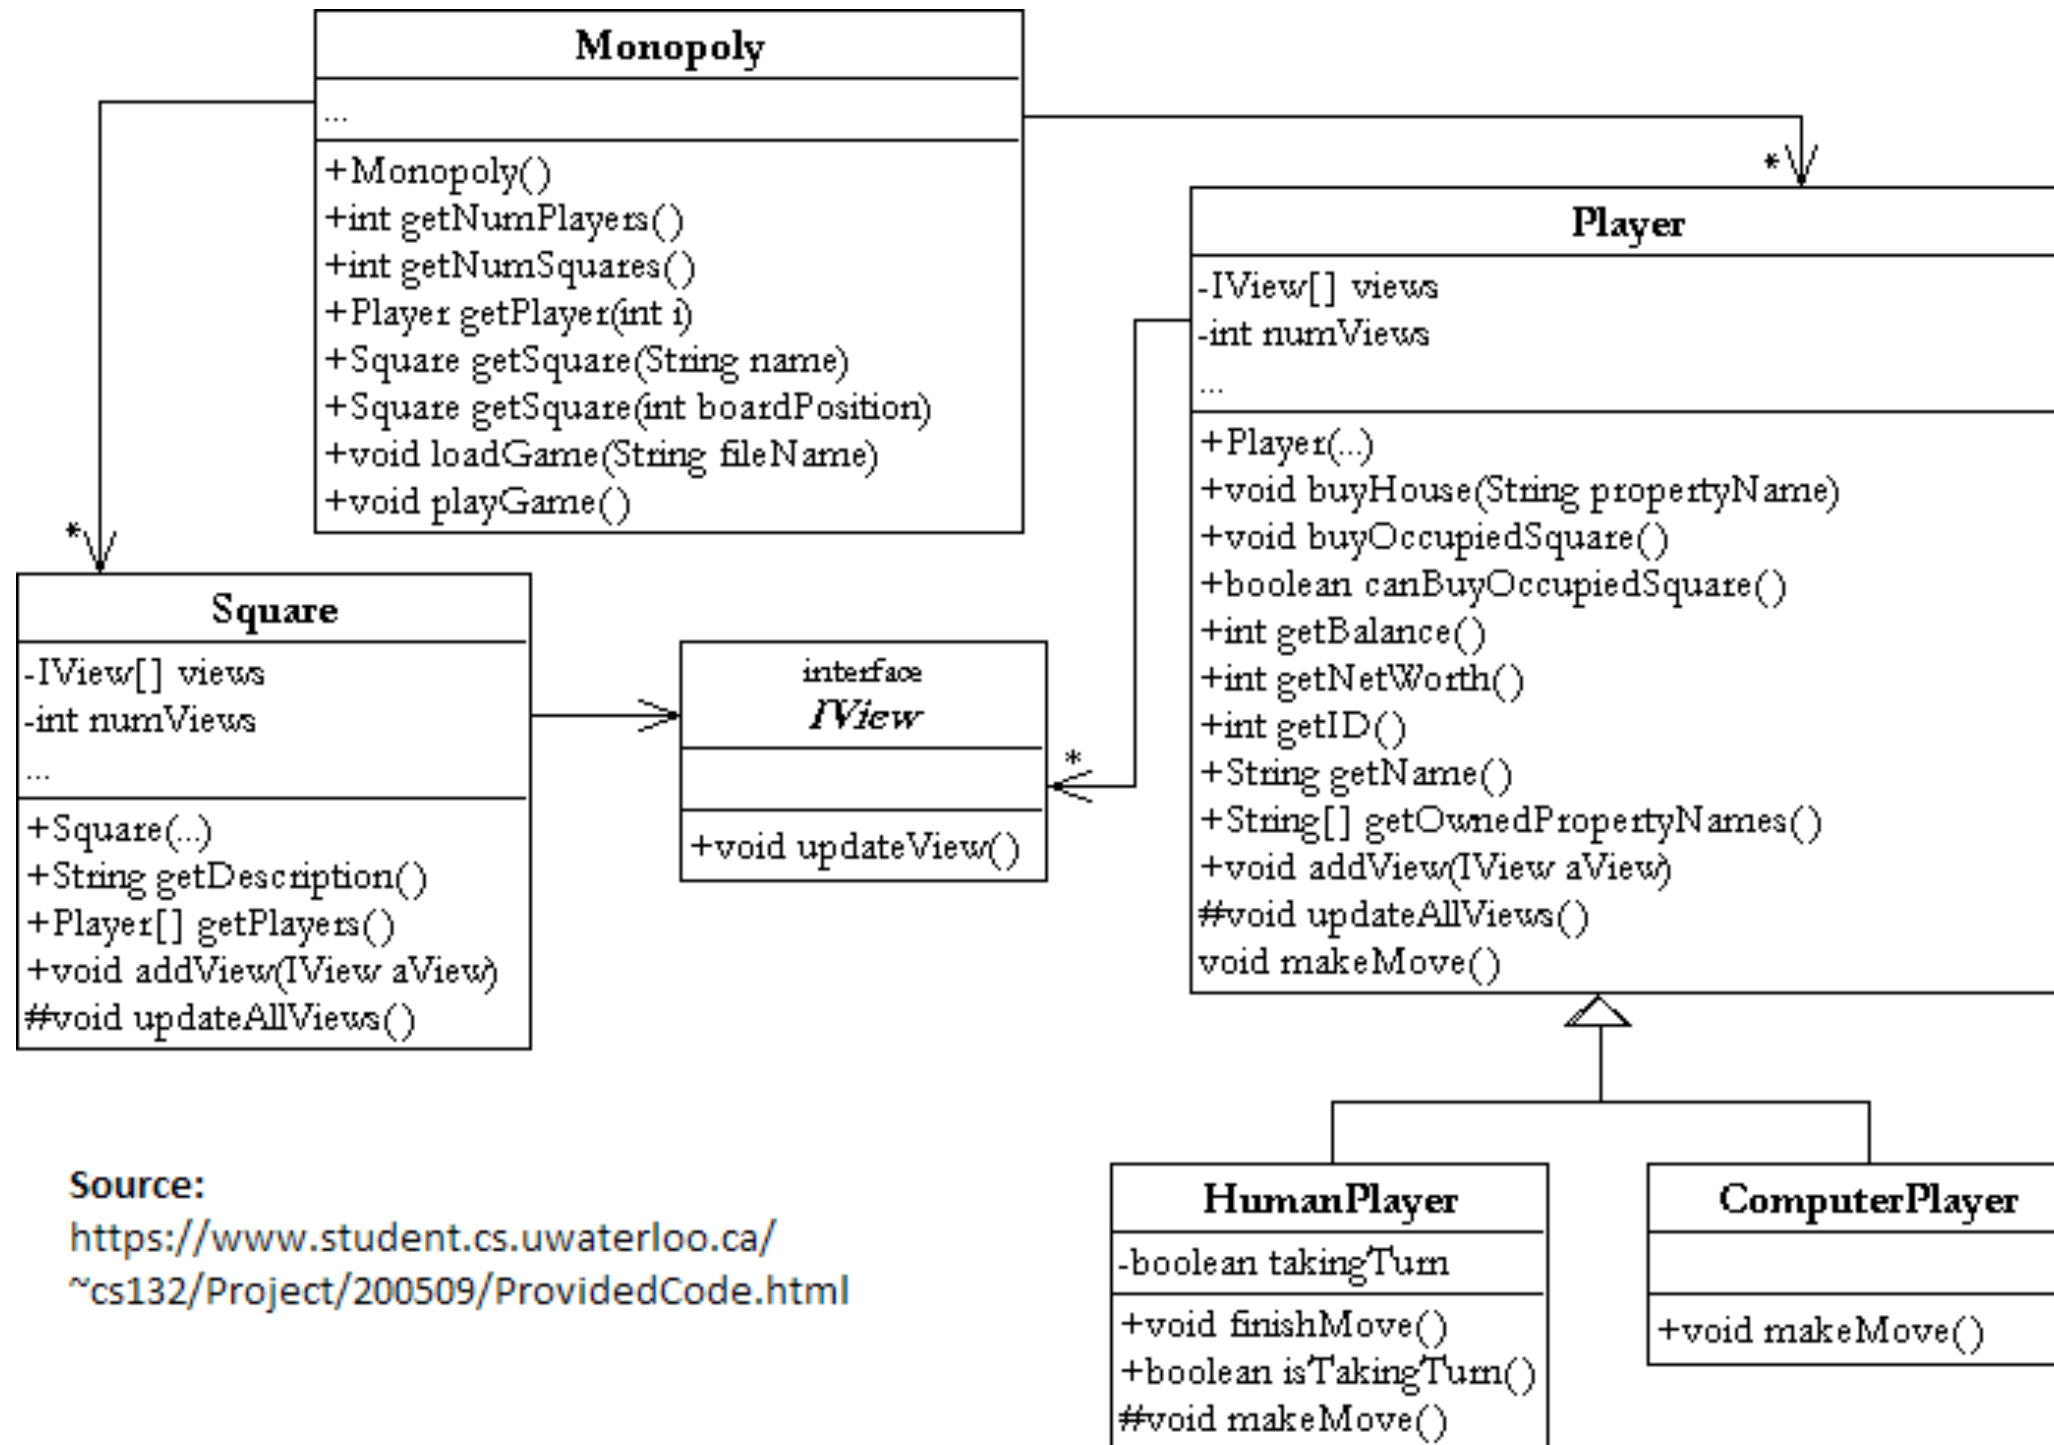
- In Handler

```
routingContext.request().bodyHandler(bodyHandler)
```

- In Body Handler

```
public void handle(Buffer buf) {  
    String json = buf.toString("UTF-8");  
    response.setStatusCode(204).end("Data saved");  
}
```

# Traditional OO Design



Source:

<https://www.student.cs.uwaterloo.ca/~cs132/Project/200509/ProvidedCode.html>

# Stateless OO Design Classes

- Model Classes - No patterns, just DTOs
- Business Classes - Design patterns apply here
- Technology Classes
- Utility Classes

# Mongodb

- Mongodb is an indexed document store.
- A document is typically something like a json structure
- Can be indexed based on any of the fields
- Can be replicated, sharded, tunable consistency with single master and leader election.

- `db.users.insert({name:'Hari',email:'hari@abc.com'})`
- `show collections`
- `db.users.update({age: { $gt: 18 }},{ $set: {status:'A'}},{multi: true})`
- `db.users.remove({status: 'D'})`

# Mongodb Queries

```
db.users.find(  
  { age: { $gt: 18 } },  
  { name: 1, address: 1 }  
) .limit(5)
```

← collection  
← query criteria  
← projection  
← cursor modifier

```
SELECT _id, name, address  
FROM users  
WHERE age > 18  
LIMIT 5
```

← projection  
← table  
← select criteria  
← cursor modifier



# Creating Index

- `db.users.createIndex( { email: 1 } )`
- `db.users.createIndex( { email: 1, name: 1 } )`
- Fully covered queries:
  - `db.users.find { email: /. *yahoo.com/},{ name: 1, _id: 0 } )`

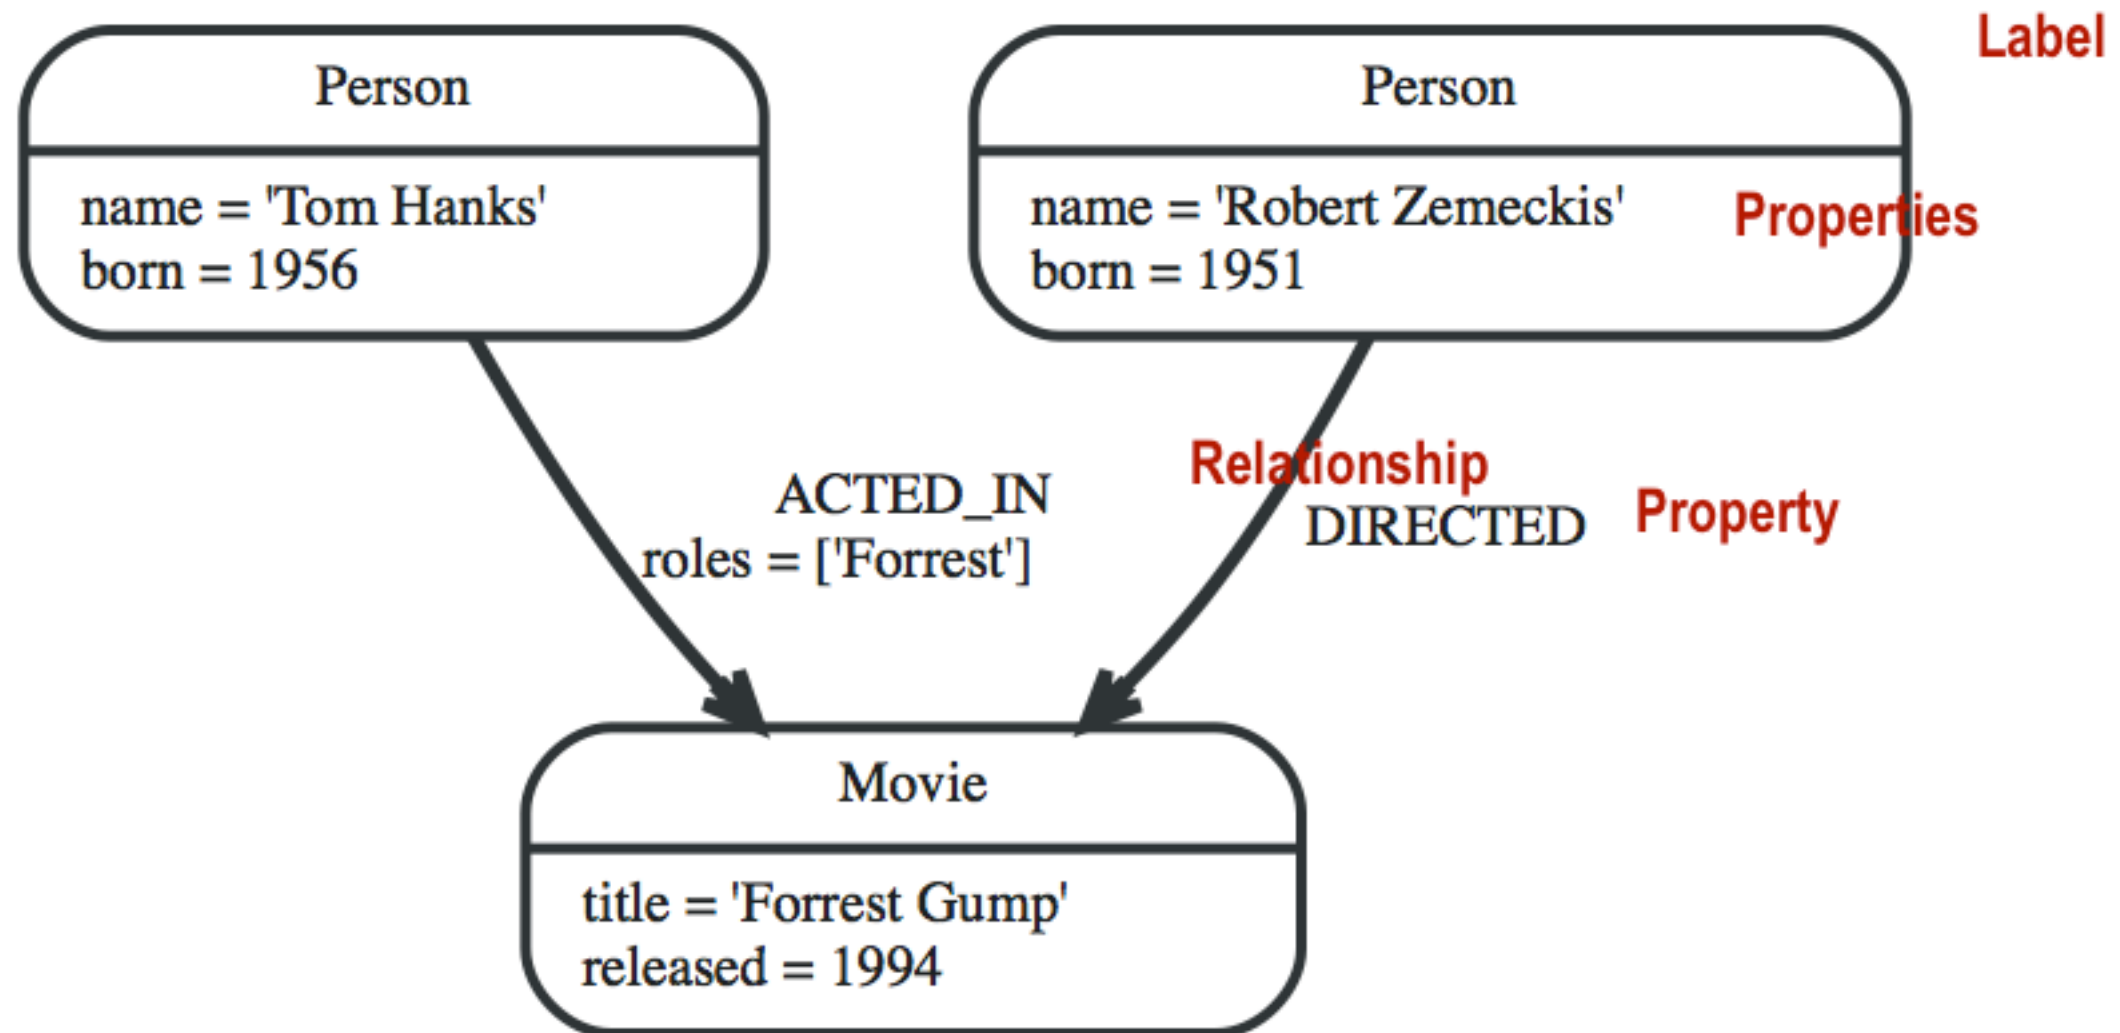
# Integrating With Java

- POM dependencies


# Neo4j

- ACID transactions,
- High availability,
- Scales to billions of nodes and relationships,
- High speed querying through traversals,
- Declarative graph query language.

# Graph




# RDBMS



# Graph


## A Property Graph



Of course.. a graph is a graph is a graph

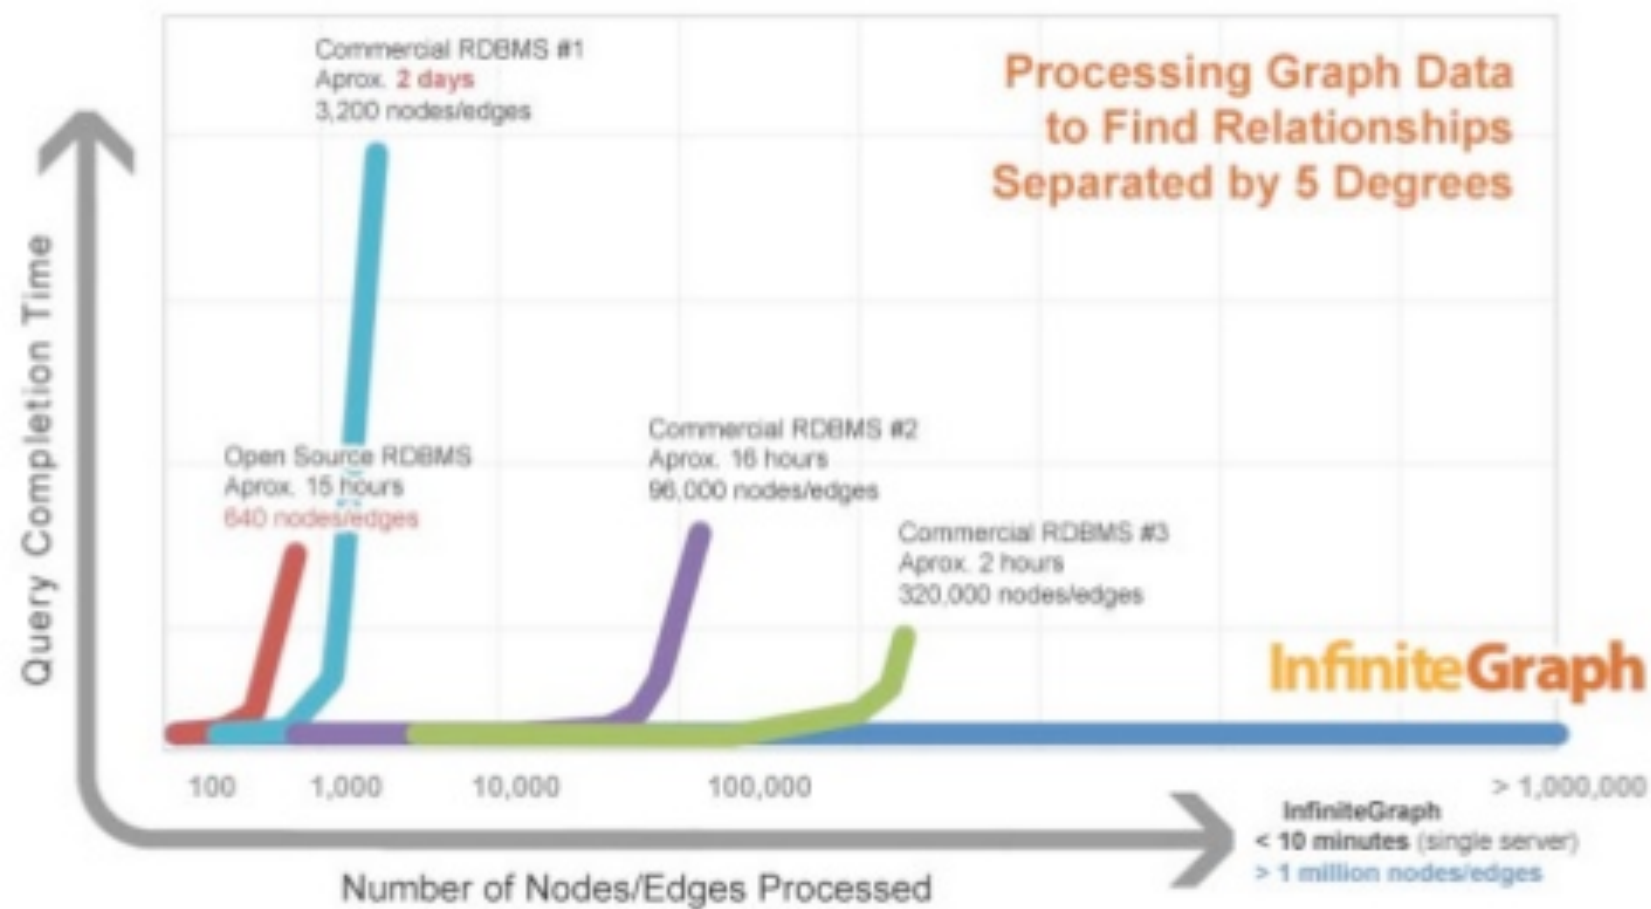


*What drugs will bind to protein X and not interact with drug Y?*





# Why use a Graph Database?





# Features

- SQL Like easy query language Neo4j CQL
- It supports Indexes by using Apache Lucence
- It supports UNIQUE constraints
- It contains a UI to execute CQL Commands : Neo4j Data Browser
- It supports full ACID
- It uses Native graph storage with Native GPE(Graph Processing Engine)
- It supports exporting of query data to JSON and XLS format

# Features

- It provides REST API for the data
- It supports two kinds of Java API: Cypher API and Native Java API to develop Java applications.

# CQL

- create (blog:Blog) - node name blog, label Blog
- create (:Blog{title:'Some blog',desc:'This is some blog'}) - Create with properties
- match(b:Blog) return b - select \* from blog b
- MATCH (b:Blog),(u:User) CREATE (b)-[r:AUTHORED\_BY ]->(u)
  - Form relationship between EVERY blog and EVERY User
- MATCH (u:User{name:'Faizal'}) CREATE (blog:Blog{title:'related blog',desc:'This is some related blog'})-[r:AUTHORED\_BY{when:'5th Sept'} ]->(u)
  - Create an object with a relationship
- match(b:Blog)-[a:AUTHORED\_BY]-(u:User{name:'Faizal'}) return a.when
  - Find all dates when Faizal authored blogs
- MATCH (b:Blog) WHERE b.title =~ '.\*Some.\*' RETURN b

# CQL

- MATCH (u:User{name:'Faizal'}) CREATE (blog:Blog{title:'related blog',desc:'This is some related blog'})-[r:AUTHORED\_BY{when:'5th Sept'} ]->(u)
  - Create an object with a relationship
- match(b:Blog)-[a:AUTHORED\_BY]-(u:User{name:'Faizal'}) return a.when
  - Find all dates when Faizal authored blogs
- MATCH (b:Blog) WHERE b.title =~ '.\*Some.\*' RETURN b

# CQL

- MATCH (b:Blog)-[rel]-(u:User) DELETE rel
  - Delete only the relationships
- MATCH (u:User{name:'Faizal'}) SET u.sex = 'Male'  
RETURN u
  - Set a field and return the node

# Java API support

- Neo4J supports two types of Java APIs
  - Native Java API
  - CQL API